



Diploma in **Data Analytics**

Lesson 4: Linear regression continued...



Linear regression continued <

Data frames <

Dates in R <

Lesson Objectives

Lesson 4



Linear regression continued

Recap linear regression

- $\hat{y} = a + bx_i$
- `lm.fit()`
- `Summary(lm.fit)`



Multiple linear regression

- More than one predictor variable
- $$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \varepsilon$$




Residual standard error

- Estimate of standard deviation of ε
- Average deviation of response from true regression line
- Measure the lack of fit of the model



R^2 statistic (RSS)


- Measures amount of variability left unexplained by model
- $0 < R^2 < 1$
- Close to 0
 - Model does not explain variability in response well
- Close to 1
 - Model explains variability in response well



Model fit statistics in R

`Summary(lm.fit)`






Data frames

Tidyverse recap

- Collection of data analytics tools contained in R for transforming and visualising data
- Tibbles
 - Data frame with tweaks




More about tibbles

(package contained in tidyverse)

- Create tibble from data frame
 - `As.tibble()`
- Create new tibble
 - `Tibble()`
- Define a tibble row by row
 - `Tribble()`





```

Terminal x Jobs x
cloud/project/
> mtcars

```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Tibble vs data frame

- Tibble have refined print method
- Shows first 10 row, columns fit on screen, large data is easy to work with
- Subsetting
- Stricter




Date in R

Lubridate package in R

- Easier to handle dates and times in R
- Built by: H. Wickham & G. Grolemund
- Maintained by: V. Spinu





Current dates and times

- Name of current timezone: `Sys.timezone()`
- Current date: `Sys.Date()`
- Current date and time: `Sys.time()`

Lubridate package

- Current date and time: `Now()`



Converting dates to strings

- Convert strings to date format:
`as.Date(y, format = "%m/%d/%Y")`
- Lubridate package:
`ymd(x)`, `mdy(x)`





Challenge >>

Create multiple linear regression models each with different predictor variables. Which model fits the data best and why?

#exploredata